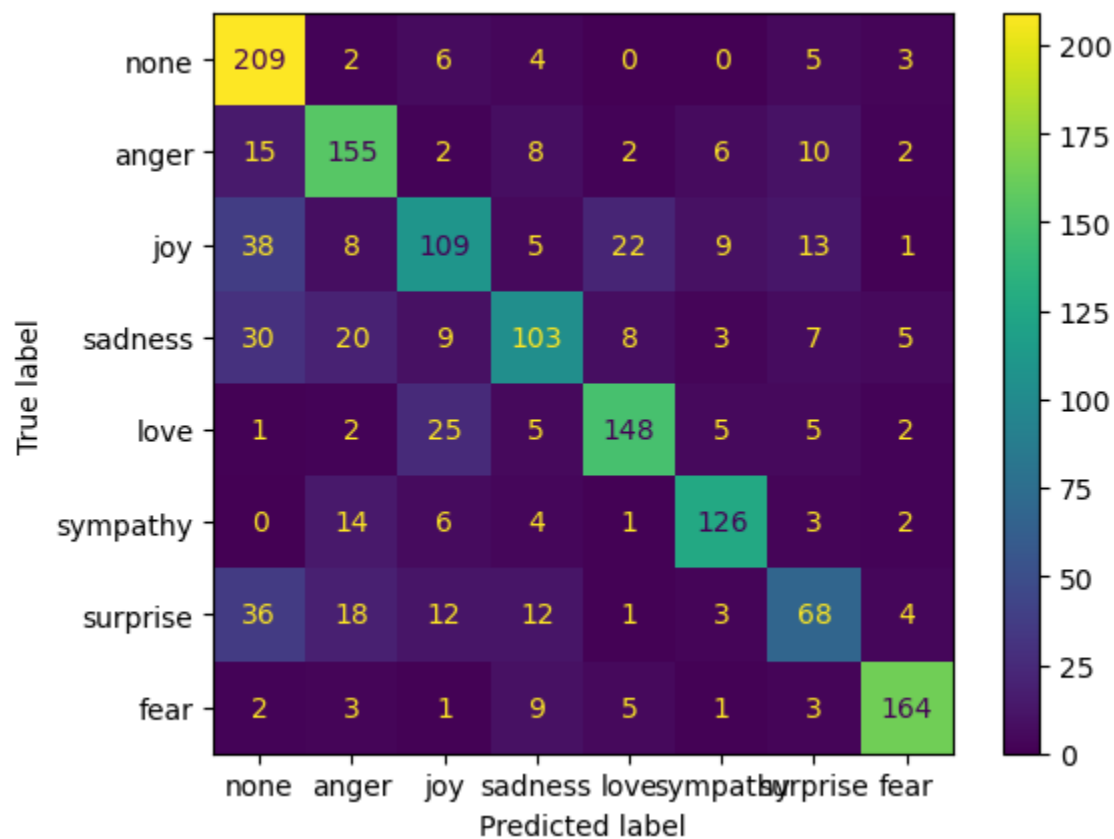
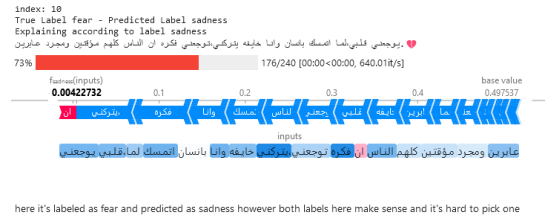


We will be using the light stemming model from now.

	precision	recall	f1-score	support
none	0.63	0.91	0.75	229
anger	0.70	0.78	0.73	200
joy	0.64	0.53	0.58	205
sadness	0.69	0.56	0.61	185
love	0.79	0.77	0.78	193
sympathy	0.82	0.81	0.82	156
surprise	0.60	0.44	0.51	154
fear	0.90	0.87	0.88	188
accuracy			0.72	1510
macro avg	0.72	0.71	0.71	1510
weighted avg	0.72	0.72	0.71	1510



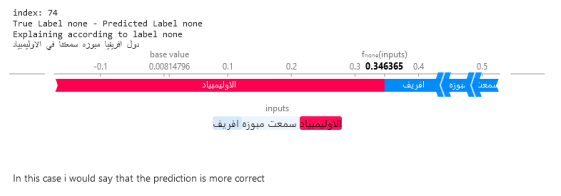
Light Stemming vs Raw Data:




Here the model made the correct prediction with light stemming.



The model stopped using stopwords as explanations.



This was a case which I agreed with the prediction more than the label, but now the model got the label correctly. The explanation is bad too. ***(try the correct word)



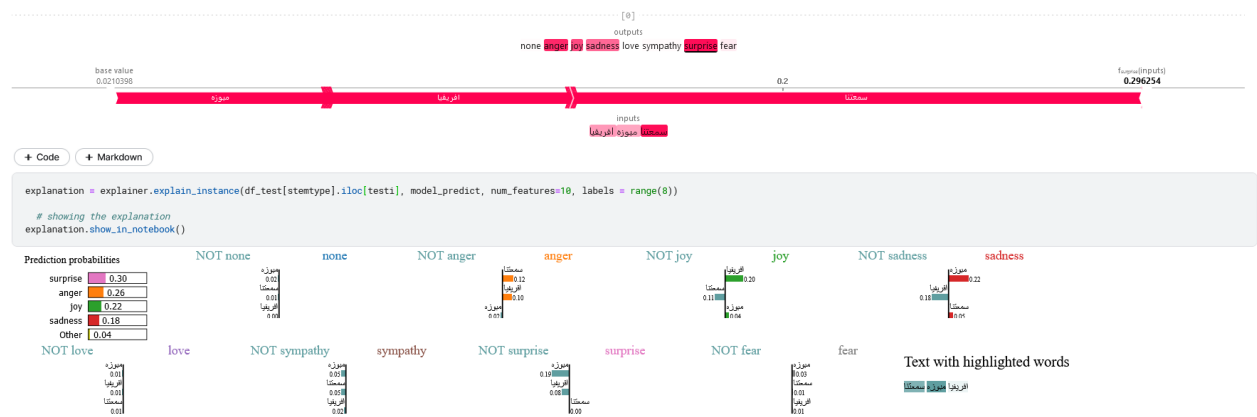
Text with highlighted words

afriqia mizre smektia alawimbiad

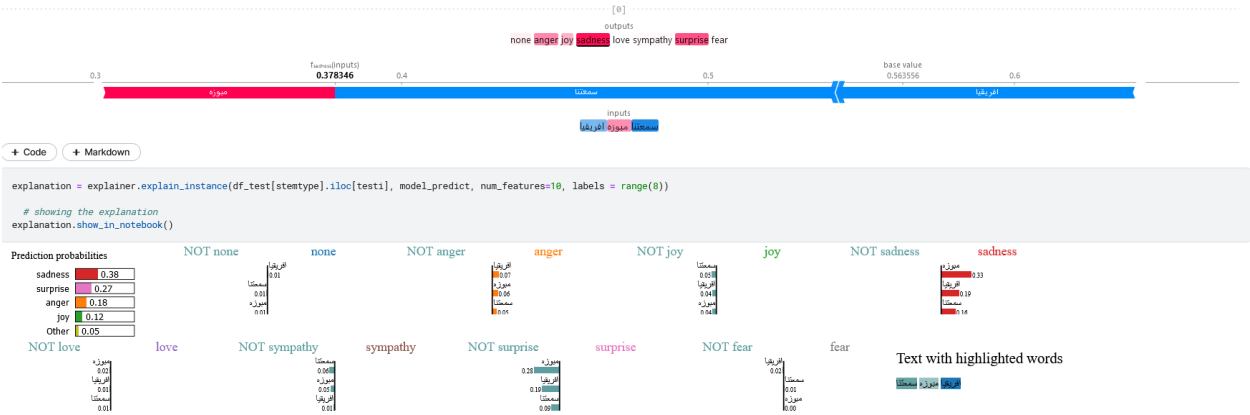
No significant change after correction?

Only that the word turned from being not-none to none but its score is not significant enough to change the outcome in either cases

Let's us try without the problematic word:




Now let us fix the word:



We returned to sadness once more.



This one just got more confusing.



The model got confused in this case, as light stemming loses some of the context, so the mistake is understandable.

Conclusion:

The explanations got better but some of the context was lost.

Shap vs Lime

Interruption:

Some errors were found in light stemming during the comparison of lime and shap.

Raw data:

💔.يوجعني قلبي،لما اتمسك بانسان وانا خايفه **يتركني،توجعني** فكره ان الناس كلهم مؤقتين ومجرد عابرين

Light Stemming:

يوجع قليل اتمسك بانس وانا خايفه **يتركني،توجع** فكره الناس مؤق ومجرد عابر حزن



While the main contributor to the label is the same, the other words are assigned different importance in different explainers.



The same words contribute to different labels according to different explainers.



Here, both have mostly the same explanation, but the words are assigned different weights.

من كثير ما احبه داني اروح احبته من ايوة بس حبيبة زينة

base value
0.0148965

f=inputs
0.450807

0.1 0.2 0.3 0.4 0.5 0.6

احبه
ايوة
حبيبة
زينة

inputs
يرفض احبته ايوة احبته كثير

Prediction probabilities

NOT love love

love	0.45
fear	0.42
anger	0.04
joy	0.04
Other	0.05

Text with highlighted words

كثير داني اروح احبته ايوة حبيبة زينة

كثير	0.25
داني	0.14
ايوة	0.11
اروح	0.11
احبته	0.07
حبيبة	0.03
زينة	0.02

Observe the weight of each words in the following examples:

[illegible]

2)

```
index: 14
True Label sadness - Predicted Label anger
Explaining according to label anger
موتدا التحكيمى الفاجر ده
```



```
explain_example(testi)
```

```
index: 14
True Label sadness - Predicted Label anger
Explaining according to label anger
موتدا التحكيمى الفاجر ده العظيم
```




3)

```
index: 15  
True Label love - Predicted Label love  
Explaining according to label love  
👉 🇲🇪 🇸🇦 🇮🇶 🇵🇰 🇯🇴 🇱🇪 🇩🇪 🇧🇭
```



```
explain_example(test1)
```

index: 15
True Label love - Predicted Label love
Explaining according to label love
👉 شئت الله



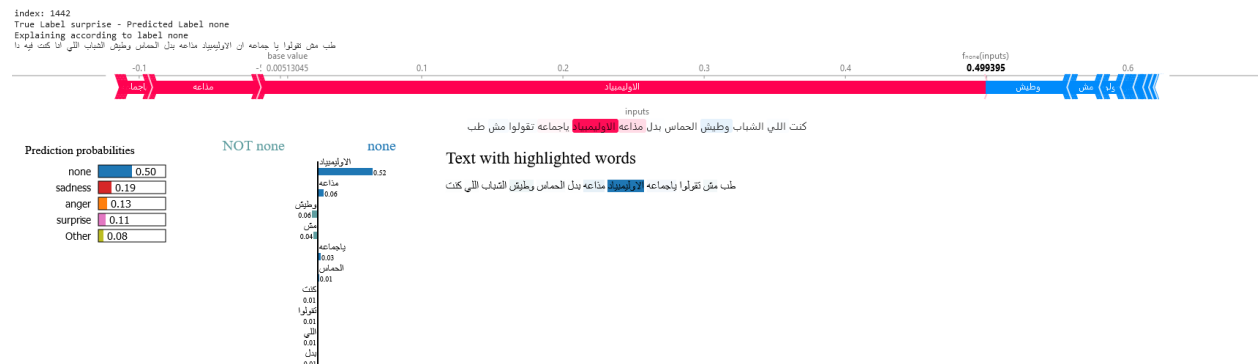
4)



Due to the nature of its calculations, SHAP is very consistent in its output. Lime on the other hand is inconsistent due to its random nature.

Explaining the same example using Lime multiple times gives different weights each time but it's usually generally correct.

The most significant word is almost always the same and the probabilities don't appear to change but it would make a difference if the values were close to each other such that a 0.01 is enough to make a difference.





As we can see here, a small change was enough to make a difference.

The last word switched from None to Not-none.

Now, we attempted to figure out why some words don't have weight in shap but have high weight in lime:



Let's return to this example once more.

The word 'معقول' has a weight of -0.09 in lime (it was verified by multiple runs) but appears to have no significant weight with shape (-0.007)

Across multiple runs I tried to remove words that has little to no weight in both lime and shap as following:

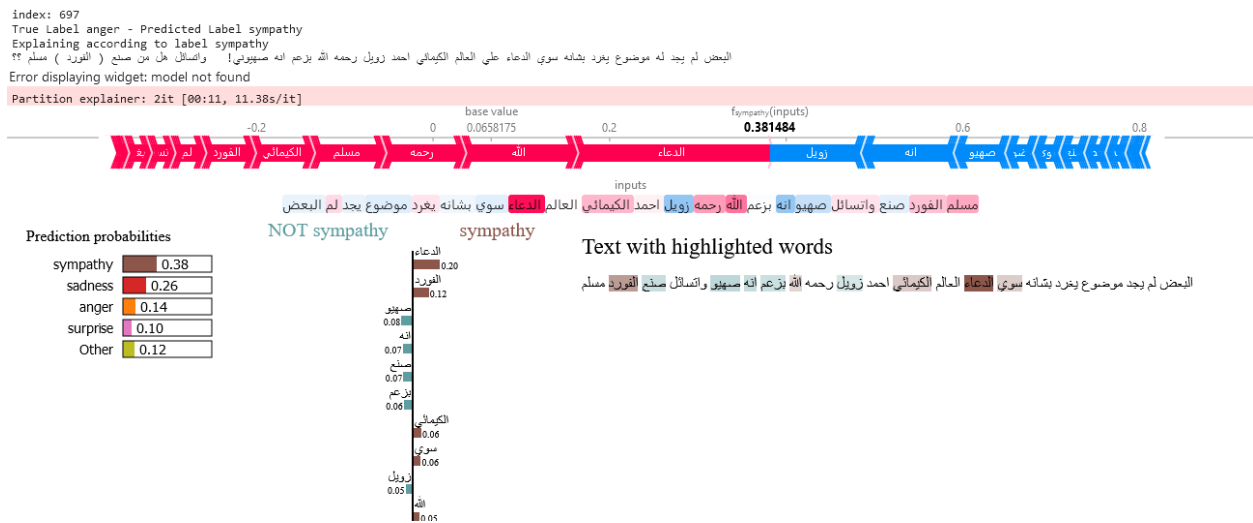




As we can observe in lime, other than the removed word, the words have the same order and a little difference in weight.

But in Shap, each removed word has a significant effect on the rest of the weights even if the removed word's weight is not significant itself.

Observe this 21-words example.



index: 697

True Label anger - Predicted Label sympathy

Explaining according to label sympathy

البعض لم يجد له موضوع يغرد بشأنه سوى الدعاء على العالم الكيماني احمد زويل رحمه الله بزم انه صهيوني! واتسائل هل من صنع (الفورد) مسلم؟؟

Error displaying widget: model not found

Partition explainer: 2it [00:11, 11.21s/it]



2)

index: 117

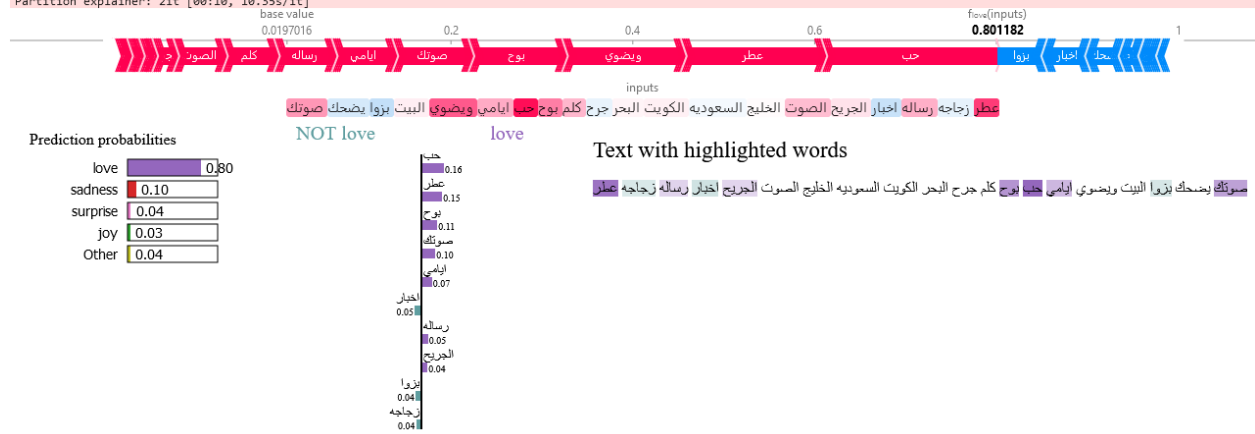
True Label love - Predicted Label love

Explaining according to label love

صوتك يضحك بزوايا البيت و يضوي ايامي ♥ بوح كلمات جرح البحرين الكويت السعودية الخليج الصوت الجريح اخبار رساله مع زجاجه عطر

Error displaying widget: model not found

Partition explainer: 2it [00:10, 10.35s/it]

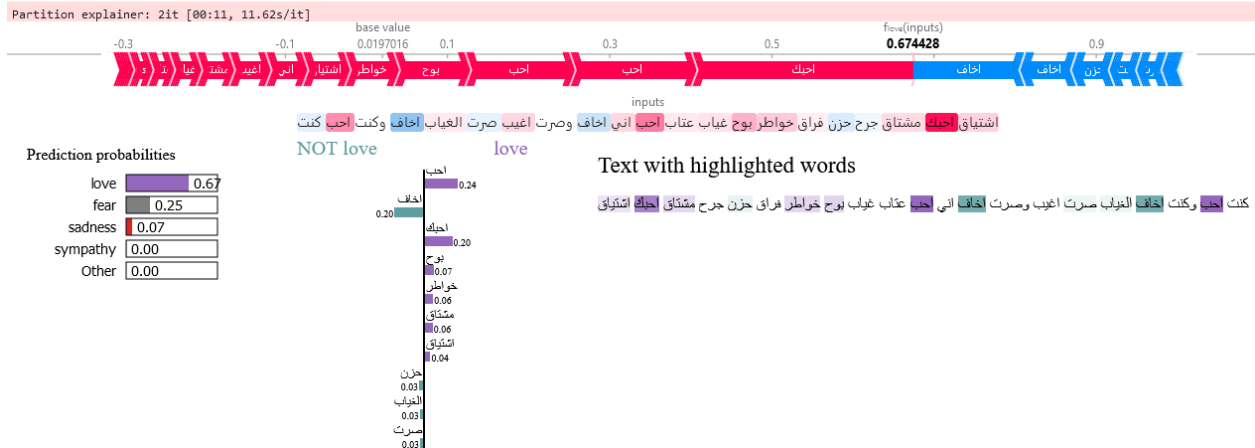


index: 117
True Label love - Predicted Label love
Explaining according to label love
صوتك يضحك بزوايا البيت و يضحكي اياي ♥ بوح كلمات جرح البحرين الكويت السعودية الخليج الصوت الجريح اخبار رساله مع زجابه عطر
Error displaying widget: model not found



3)

index: 838
True Label love - Predicted Label love
Explaining according to label love
كنت احب وكنت احاف من الغياب صرت اعيب وصرت احاف اني احب عتاب غياب بوح خواطر فراق حزن وله جرح مشتاق احبك اشتياق
Error displaying widget: model not found

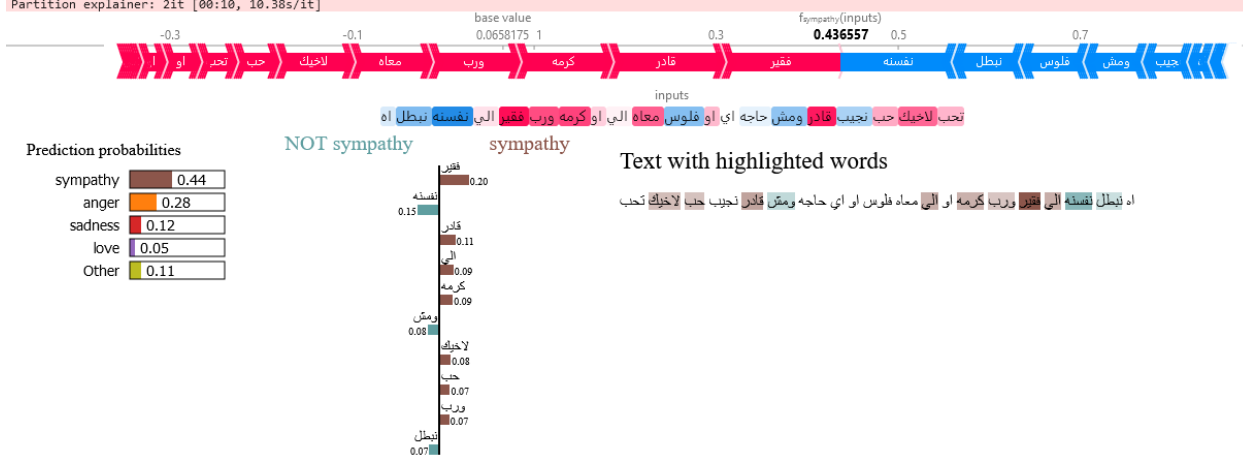


index: 838
True Label love - Predicted Label love
Explaining according to label love
كنت احب وكنت احاف من الغياب صرت اعيب وصرت احاف اني احب عتاب غياب بوح خواطر فراق حزن وله جرح مشتاق احبك اشتياق
Error displaying widget: model not found

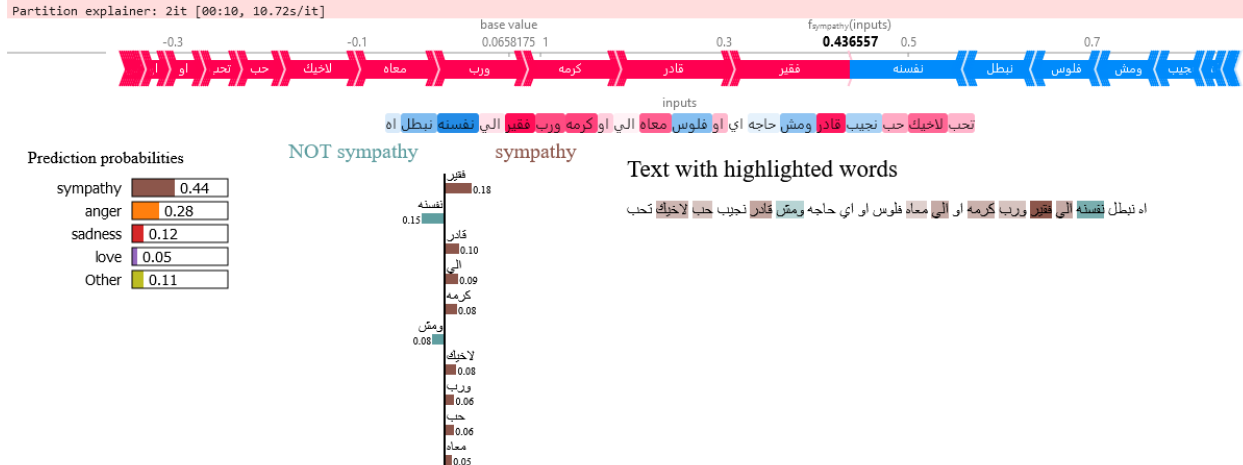


4)

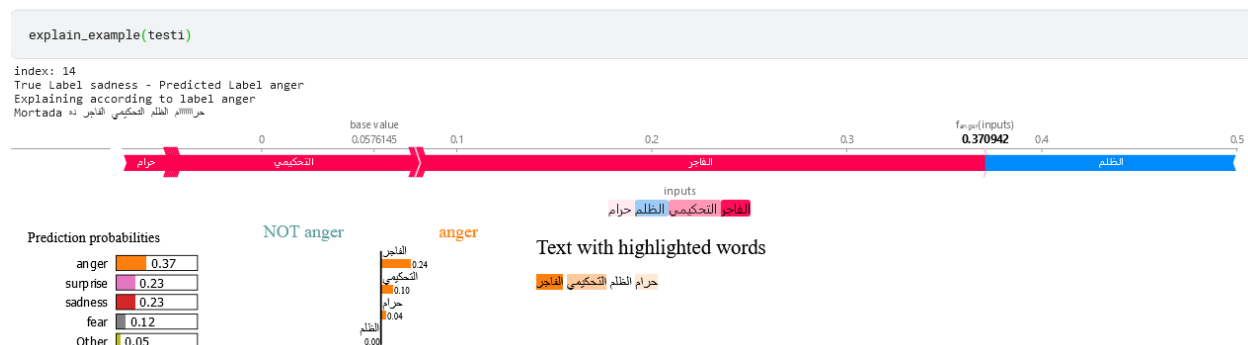
```
index: 19
True Label sympathy - Predicted Label sympathy
Explaining according to label sympathy
اډ لو نېټل نفسه من الي كان فقير وربنا كرمه او من الي معاه فلوس او اي حاجه نفسا فيها ومش قاترين نجيبا حب لايخيك ما تحب نفسك
Error displaying widget: model not found
Partition explainer: 2it [00:10, 10.38s/it]
```



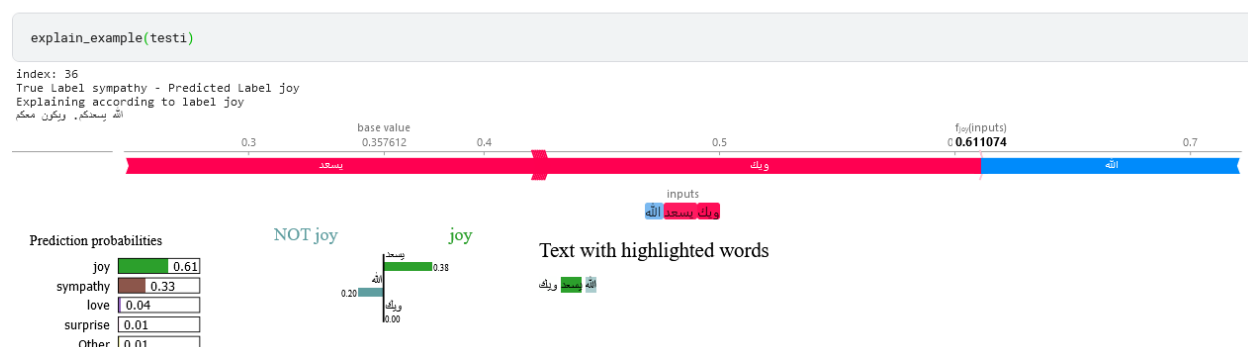

```
index: 19
True Label sympathy - Predicted Label sympathy
Explaining according to label sympathy
اډ لو نېټل نفسه من الي كان فقير وربنا كرمه او من الي معاه فلوس او اي حاجه نفسا فيها ومش قاترين نجيبا حب لايخيك ما تحب نفسك
Error displaying widget: model not found
Partition explainer: 2it [00:10, 10.72s/it]
```



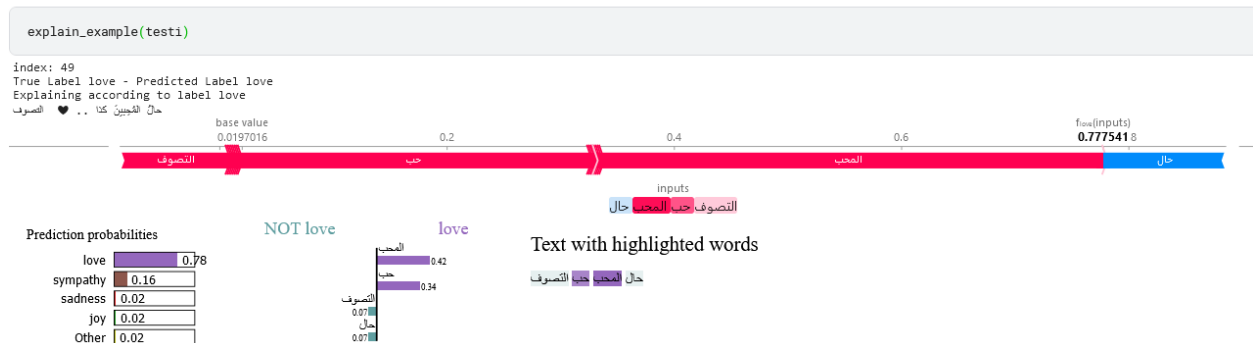
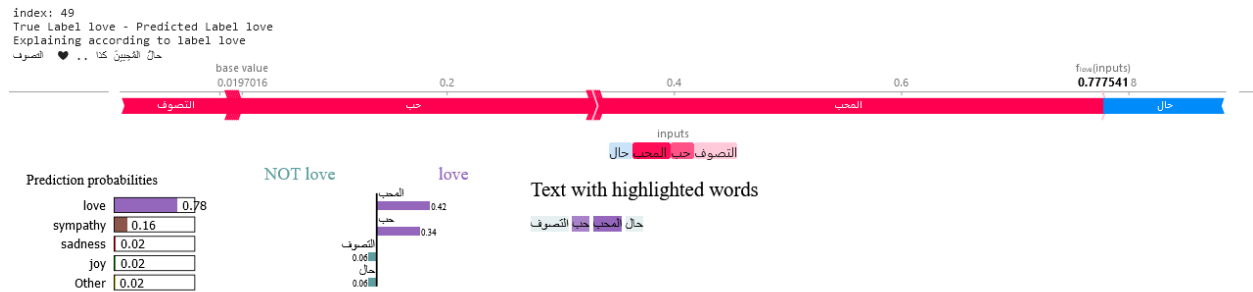
Let's observe small examples now:



2)



3)



4)



It seems that lime is unable to explain it properly and It may be because Lime generates 5000 random samples and while those examples are enough to cover most sentences, it appears that it is not enough to cover all permutations possible in a long sentence so, the changes between different runs are a bit higher.

The highest change in short examples tend to be 0.01 but in long sentences it can reach 0.02 of change in multiple words at once.

Conclusion:

Due to the way of calculation, Lime has little in the way of stability when running the same example multiple times but is surprisingly robust when some of the input is changed or dropped.

Shap on the other hand is very consistent in its calculations when running the same example, but it is greatly affected by small changes in the input and seemingly unrelated words may have significant changes in the output.

Interruption (again):

Some errors were found in light stemming during the comparison of lime and shap.

Raw data:

دول **افريقيا** ميوزه سمعتنا في الاولمبياد

Light Stemming:

افريف ميوزه سمعت الاولمبياد

There is a spelling mistake here, and some removed stopwords that need to be looked into.

And another one:

Raw data:


محدث عارف هما تعبوا اد ايه عشان يوصلوا الاولمبياد و هم مش ناقصين تعليقات الجهله في اللعبه لادائهم و لا تعليقات **المصريين** **عموما** علي لبسهم

Light Stemming:

محدث عارف تعب اد عشان يوصل الاولمبياد مش ناقص تعليق الجهله اللعبه لاداء تعليق **المصر** **عمو** لبس

There are some removed stopwords that need to be looked into.

Using Cleaned data without stemming:




There was this error that I attributed to a loss of context but without stemming, the model got it correctly once more:



So, context can be lost with stemming but it appears that without stemming, it may get better.

Since we will be using cleaned data now, here are the metrics:

	precision	recall	f1-score	support
none	0.64	0.89	0.74	229
anger	0.71	0.79	0.75	200
joy	0.62	0.56	0.58	205
sadness	0.68	0.55	0.61	185
love	0.79	0.76	0.77	193
sympathy	0.85	0.83	0.84	156
surprise	0.59	0.44	0.50	154
fear	0.91	0.87	0.89	188
accuracy			0.72	1510
macro avg	0.72	0.71	0.71	1510
weighted avg	0.72	0.72	0.71	1510



By comparing the confusion matrix with the previous one, we can observe that the box that indicates (true = none, pred = joy) increased from 6 to 10.

That was despite the fact that context appears to be more complete without stemming so let's examine those 10 samples.

76% 160/210 [00:00<00:00, 738.97%/s]

base value 0.248209


$f_{\text{in}}(\text{inputs})$ 0.463062

0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8

ساز در و جوار یاص تریکه فرج وحشتنا یوصل عیاره الومیدار الکاشن اعم

inputs

فرج باعم وحشتنا الکاشن وشاره البرازیل وچول باحول صلاح تریکه یاص عیاره کانت الومیدار



Explaining according to Layer 20

شود به کمک این لایه نتایج حاصل از مدل را می توان توضیح داد

0.1 0.2 0.3 0.4 0.5 0.6

base value 0.248209

0.497619 $f_{20}(inputs)$

یخچالی رانکویت احمد شبنو عیسیا زینبا



Explaining according to label joy
بار، بارهوا اسم مصدر بانهما راجعاً لهما شجوا مصر في الاوليين

base value 0.248209

f_{net}(inputs) 0.479533

0.2 0.3 0.4

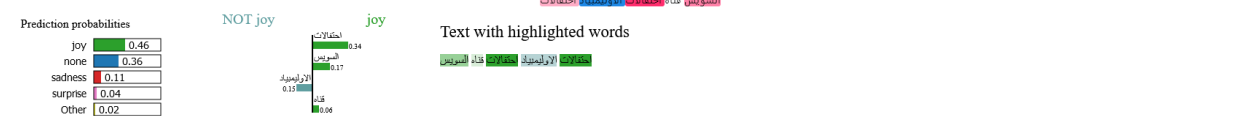
راجعهها راجعهم اسم رافعا رافعها رافعهم اسم رافعا رافعهم



الاحتمالات الأولية π_0 (inputs) base value 0.248209

0 0.1 0.2 0.3 0.4 0.5 0.6 0.7

الاحتمالات المتوسطة الاحتمالات الأولية



البيانات: 23 و 22 دب. - الإصدار في تاريخ الإصدار والإصدار الرئيسي في التاريخ حرج التمس 28

Category	Value
التمس	0.05988
حراج	0.248209
حراج	0.45988
الأوراسيا	0.6
التاريخ	0.8
حراج	0.8




Figure 1 illustrates the sentiment analysis pipeline. The top part shows the base value (0.248209) and the final output (0.673445) on a scale from 0.2 to 0.7. The input text is "مع فرح الأولمبياد الخامس المركز واحد رقم مصر". The bottom part shows the prediction probabilities for various sentiment classes: joy (0.67), none (0.19), surprise (0.11), sadness (0.02), and Other (0.01). The right part shows the distribution of sentiment classes across the input words, with "فرح" having the highest probability (0.39).

Figure 1 illustrates the process of sentiment classification using a neural network. The top part shows a sequence of words: "التحسيس", "أوى", "ياغفري", "جانو", "جانو", "الأولمقياد", "ياخدع", "لايحب", "ميداليات". Below this, a bar chart shows the prediction probabilities for each word: joy (0.62), surprise (0.19), none (0.14), anger (0.03), and Other (0.02). The bottom part shows the text "ادرج الأولمقياد ميداليات ييجوب ياخدع أوى ياخدع ميداليات الأولمقياد ادرج التحسيس" with the words "ياغفري", "أوى", "ياخدع", and "التحسيس" highlighted in green.

base value
0.246209

F1 (Inputs)
0.450508

عزلة سودا التشاؤم المحيرة الوليمية اما اولادنا عزلة

عزلة اولادنا والاولاد حقتنا الاوليمية اما سودا العزلة

Prediction probabilities

joy	0.45
none	0.31
sympathy	0.12
sadness	0.07
Other	0.05

Text with highlighted words

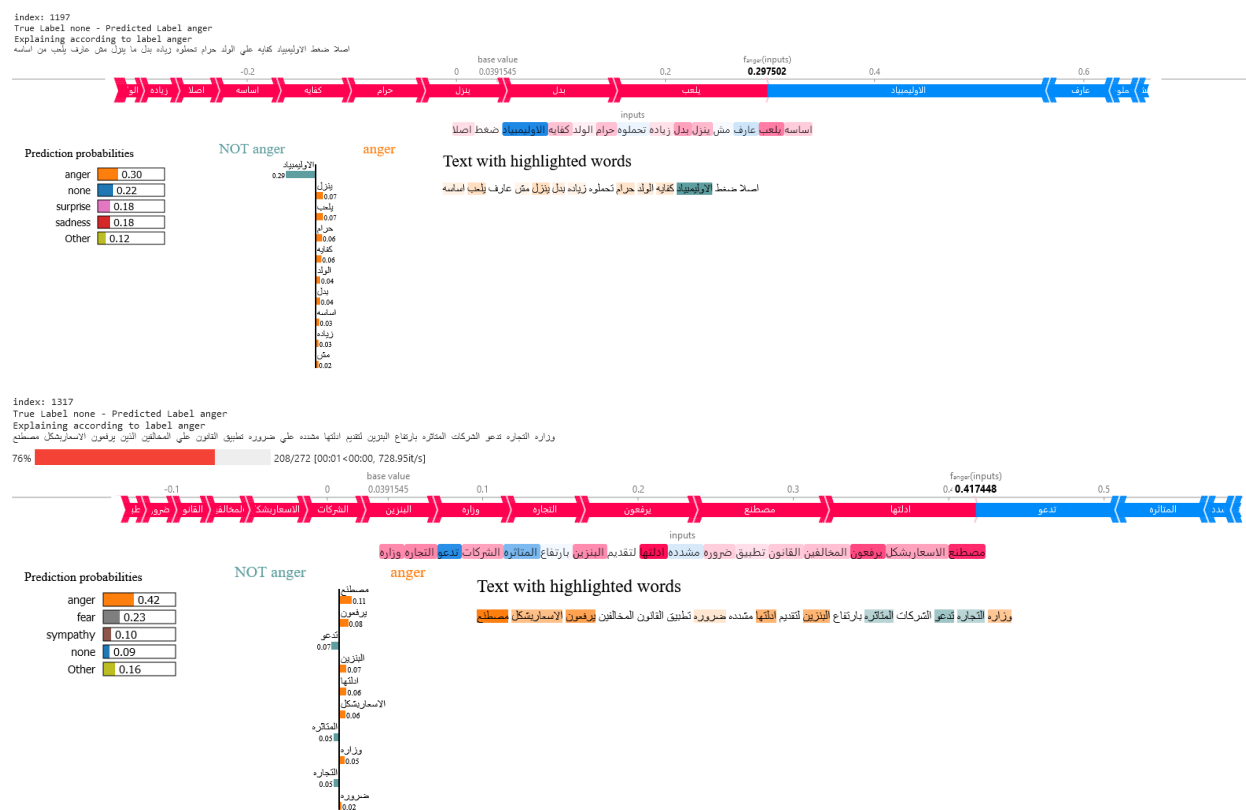
عزلة اولادنا والاولاد حقتنا الاوليمية اما سودا العزلة

Some examples are understandable why the model predicted them as joy, while others are just confusing to me and the whole context did not help either.

The just context made the model more confused.

So, let's examine something a bit less confusing.

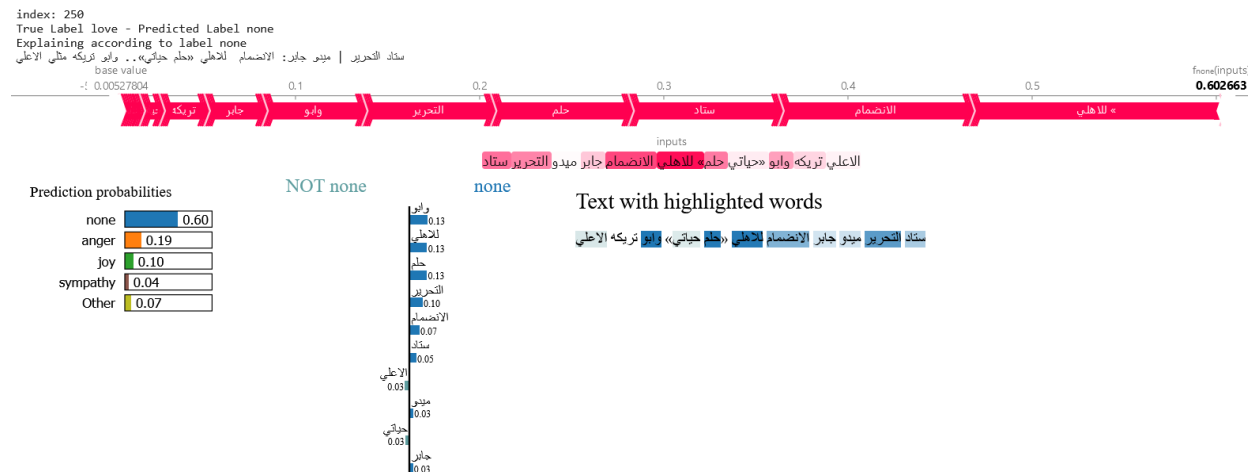
(true = none, pred = anger, total error = 3)





The first and last example are understandable why the model would label them as anger, and I am inclined to agree with its explanation. Especially the last one, as 3 curse words in one sentence is a convincing argument for why it is predicted anger.

(true = love, pred = none, total error = 1)



Most of the love samples in the dataset are the romantic kind of love, so the problem here is a lack of data in this very specific area.

(true = love, pred = joy, total error = 30)

index: 203

True Label love - Predicted Label joy

Original: احلام العنزي فخر الكويت الف الف مبرووووك بنت قبيلتي وبنت ديرتبيبي????????????????
الديره كل يوم اسم رافع اسمنا بشاره خير ان شاء الله

Cleaned: احلام العنزي فخر الكويت الف الف مبروك بنت قبيلتي وبنت ديرتي كفو عيال الديره يوم اسم رافع اسمنا بشاره خير شاء الله

I would say this one is mixed feelings.

index: 594

True Label love - Predicted Label joy

Original: : الدوله العثمانيه الاسلاميه ستبقي رغم انوف الاعداء ربي احفظ تركيا المسلمه من مكر الماكرين واحفظ كل بلاد
المسلمين ♥

Cleaned: الدوله العثمانيه الاسلاميه ستبقي رغم انوف الاعداء ربي احفظ تركيا المسلمه مكر الماكرين واحفظ بلاد المسلمين

This one supports my earlier argument about romantic love and love for one's country.

index: 823

True Label love - Predicted Label joy

Original: لون حياتك بحب وطاعه الله فهي طويله قصيره ... ايمان تقوه

Cleaned: لون حياتك بحب وطاعه الله طويله قصيره ايمان تقوه

Same as the above.

Conclusion on True = love:

It appears that most of them are predicted wrongly due to a few reasons:

1. The difference between romantic and non-romantic love
2. Mixed feelings
3. Misunderstood context
4. Generally confusing such as the following example:

index: 1430

True Label love - Predicted Label joy

Original: الهلال الساعه 7:30 بالتوفيق شباب الاتحاد & الاربعاء ديربي.. اتحاد الشرس

Cleaned: الاربعاء ديربي اتحاد الشرس الهلال الساعه بالتوفيق شباب الاتحاد

(true = surprise, pred = none, total error = 35)

All of them are due to “الاولمبياد”

ALL OF THEM!!!